

Generating text with reinforcement learning

Maximilian Balthasar Mansky

October 25, 2020

Goal of this project

Machine learning approaches to text generation try to replicate a given text, for example by predicting the next word for a sentence. While the results can be interesting, this approach has obvious limitations.¹ Text can also be generated from structured information, such as information about somebody's life.² These approaches have the drawback that the machines only learn to imitate, in the sense that it reduces the distance between the generated text and the truth. This well-put in an article by Karpathy³ and its response⁴. In this project, I want to explore whether a reinforcement learning approach provides different results.

Instead of attempting to generate text by imitation, we can check a generated solution for its correctness. This is done by implementing grammar, vocabulary and further checks in the environments that the machine interacts with. As an example, part of the environment might be responsible for evaluating word order, another for measuring word variations. Text is useful because it is easily verified by humans, compared to game solving strategies. It also allows for a wide variety of possible solutions. "The quick brown fox jumps over the lazy dog" is equally grammatically valid as "the lazy dog is jumped over by the quick brown fox".

Basic structure

There are two parts to reinforcement learning, the agent and an environment. The agent interacts with the environment and receives feedback from it. Conversely, the environment receives input from the agent to change and evaluates the new state. The result of the evaluation is passed back to the agent and forms the basis for its learning.

A simple example is a state environment. Consider a five-state environment, consisting of a negative reward state, a positive reward state, two intermediary and a starting state.

The agent then chooses its actions based on its current state. Possible actions are left or right (It is left to the environment to decide whether a state change is valid. Going left on the left-most cell is often interpreted as not changing state.) and the agent is given information about the result of that change (nil for changing to one of the center states, ± 1 for the outermost ones). How the agent learns is dependant on its implementation (i.e. whether it keeps a history, how the initial interaction is treated, how learning is implemented...). The environment is only responsible for evaluating the actions, the agent is responsible for learning from the feedback.

¹Jey Han Lau, Trevor Cohn, Timothy Baldwin, and Adam Hammond. This ai poet mastered rhythm, rhyme, and natural language to write like shakespeare, 2020. URL [ThisAIPoetMasteredRhythm,Rhyme,andNaturalLanguagetoWriteLikeShakespeare](#)

²Rémi Lebre, David Grangier, and Michael Auli. Neural text generation from structured data with application to the biography domain, 2010. URL <https://research.fb.com/wp-content/uploads/2017/02/neural-text-generation-emnlp-camera-ready.pdf>

³Andrej Karpathy. The unreasonable effectiveness of recurrent neural networks, 2015. URL <http://karpathy.github.io/2015/05/21/rnn-effectiveness>

⁴Yoav Goldberg. The unreasonable effectiveness of character-level language models, 2015. URL <https://nbviewer.jupyter.org/gist/yoavg/d76121dfde2618422139>

-1		start		+1
----	--	-------	--	----

Figure 1: A simplified example of an environment. The agent starts in the center cell and can move left or right. The left-most cell contains a negative reward (denoted by -1), the rightmost one a positive reward ($+1$).

Environment

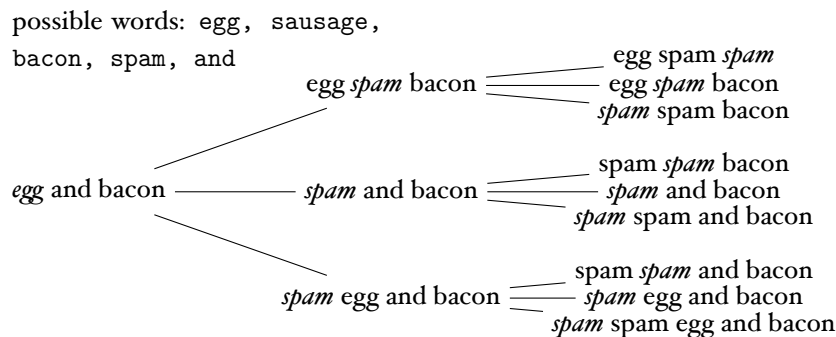
In this project, the Environment needs to check whether the current state (a text or sentence to be built by the agent) is correct and assign some form of feedback based on it. Each item of correctness will be implemented separately in grammar model. Possible models are:

- Sentence length
- Presence of nouns, adjectives, verbs
- Position of words
- Word order/grammatical correctness
- Sentence variation

As current agent learning frameworks can only work with a single reward dimension, the model output must be coerced into a single output by the Environment. It may be worthwhile to explore whether gradually returning the reward from more complex grammar models⁵ is useful in guiding the agent. For evaluating the states, the Environment must at least keep track of the current sentence. In this context, a state means a particular sentence, not the individual words.

Within a state, words itself need to be encoded somehow. Either by using a one-hot encoding (limited to the most common words) or a small word2vec model based on a text. The latter has the advantage that the size of the vocabulary can be smoothly increased, by adding more words near to similar ones. In the latter case of a word2vec encoding, it probably also makes sense to include some sort of word accuracy reward, indicating how close the chosen word vector is to the vector output of the agent.

The actions changing states can be implemented in several ways. Either by generating a new state from scratch (that is, the agent generates a sentence for evaluation) or by manipulating the existing state, for example by changing the current word or the ones next to it. An action would consist of two parts: Deciding which to manipulate (current word, previous or next one or quit/finish sentence⁶) and choosing which word to use.



Moving between states can give different paths to the same state (for example "spam spam bacon"), but the reward derived from it should be the same.⁷

⁵ Implementation idea: Use some form of weighted average from the outputs. In the beginning the model may focus on producing sentences of the right length and with a bit of variation (80% sentence length, 20% word presence) and as learning progresses, focus more on the higher end of complexity (20% word position, 30% word order, 50% sentence variation).

⁶ Maybe it is worthwhile to also include insertion of words? Or just move, do not change word?

Figure 2: Example of state transfer from an initial sentence, with the possible words replacements indicated by *emphasized* text.

⁷ The Environment is ignorant of its previous states and keeps no memory.

Agent

The Agent is a program that interacts with the environment and learns from it. The Agent receives information from the Environment – The current sentence and its associated quality in the form a reward. By building experience through interaction, the Agent can build a model of the environment and move towards higher rewards.

In order to learn from the Environment, the Agent needs two things: Interaction, and memory of the best actions. Initially the Agent is clueless about its surrounding and needs to explore. Since there is no information available, the best action is to explore randomly. The interactions and their results are stored and as a random string of interactions leads to positive results, the machine starts to learn from its environment.⁸

⁸ Andy Thomas. Reinforcement learning tutorial with tensorflow, 2018. URL <https://adventuresinmachinelearning.com/reinforcement-learning-tensorflow/>

Implementation

Basic interaction

A first version of the grammar environment is implemented in `grammar_v1.py`. The agent interacts through `.step(action)`. The environment assumes that `action` is composed of a choice and a word vector. The choice is a one-hot encoding of the following options:

Version 1

- o Replace the word vector at the current position
- 1 Move the cursor position towards the negative without changing the word vector, to a maximum of -1 .
- 2 Move the cursor position towards the positive without changing the word vector, to a maximum of `current word length + 1`.
- 3 End the current interaction with the environment.

There is no check on the length of the one-hot vector, meaning that the last option (3) can be ignored. The encoding is read with an `argmax`, the values do not have to sum up to 1. Indices above 3 lead to an `NotImplementedError`.

References

- Yoav Goldberg. The unreasonable effectiveness of character-level language models, 2015. URL <https://nbviewer.jupyter.org/gist/yoavg/d76121dfde2618422139>.
- Andrej Karpathy. The unreasonable effectiveness of recurrent neural networks, 2015. URL <http://karpathy.github.io/2015/05/21/rnn-effectiveness>.
- Jey Han Lau, Trevor Cohn, Timothy Baldwin, and Adam Hammond. This ai poet mastered rhythm, rhyme, and natural language to write like shakespeare, 2020. URL [ThisAIPoetMasteredRhythm,Rhyme,andNaturalLanguageToWriteLikeShakespeare](#).
- Rémi Lebre, David Grangier, and Michael Auli. Neural text generation from structured data with application to the biography domain, 2010.

URL <https://research.fb.com/wp-content/uploads/2017/02/neural-text-generation-emnlp-camera-ready.pdf>.

Andy Thomas. Reinforcement learning tutorial with tensorflow, 2018. URL <https://adventuresinmachinelearning.com/reinforcement-learning-tensorflow/>.